
Name: Victory Chukwudi
Course: Statistical Machine Learning

Assignment Number: 2
Date: December 16, 2024

Introduction

The prostate cancer dataset contains DNA microarray gene expression data, with observations corresponding to both cancerous and non-cancerous subjects. This dataset has been widely used to evaluate machine learning algorithms for classification tasks based on gene expression levels. The aim of this analysis is to assess the dataset in terms of its shape, the distribution of the response variable, the identification of the most significant predictors using the Kruskal-Wallis test, and the visualization of these findings to facilitate model interpretation. Specifically, we compare the performance of various learning machines (1NN, 7NN, 9NN, Trees with different cp values) based on their test error rates, but first, we focus on understanding the characteristics of the dataset.

1. Dataset Shape and Structure

Sample Size and Dimensionality

The dataset consists of 79 samples and 500 predictor variables (X1 to X500), with one categorical response variable (Y). The dimensions of this dataset suggest a high-dimensional setting, where the number of features (predictors) far exceeds the number of samples. This characteristic typically poses a risk of overfitting, where models may perform well on training data but fail to generalize effectively to unseen data. Furthermore, in such high-dimensional spaces, models may require dimensionality reduction techniques (e.g., Principal Component Analysis or feature selection) to improve robustness and reduce model complexity.

- **Sample Size (n):** 79
- **Dimensionality (p):** 500
- **Predictor-to-Sample Ratio (p/n):** 6.33
- **Missing Values:** The dataset does not contain missing values, which is advantageous for model training and analysis, eliminating the need for imputation.

Statistical Overview of the Data

- **Response Variable (Y):** Categorical (binary) variable with two levels: 0 (non-cancer) and 1 (cancer).
- **Predictor Variables (X1 to X500):** Continuous variables representing gene expression levels for various genes.
- **Data Type Summary:** The response is binary, while all predictor variables are continuous, indicating that this is a classical binary classification problem with high-dimensional predictors.

2. Distribution of the Response Variable

A critical step in analyzing any dataset is understanding the distribution of the response variable. The distribution of the response (Y) is visualized using a barplot.




Figure 1: Distribution of Response Variable (Y)

The counts of each class are as follows:

- **Cancer (Y=1):** 42 samples
- **Non-cancer (Y=0):** 37 samples

From the barplot, it is evident that the dataset is relatively balanced in terms of the response classes, with a slight tilt towards more cancer cases (Y=1). The balanced distribution of the response variable helps in building machine learning models without significant class imbalance, which can otherwise lead to biased models or inaccurate predictions.

3. Identification of the Most Powerful Predictors

To assess which gene expression levels most effectively differentiate between the two groups (cancer vs. non-cancer), we applied the Kruskal-Wallis test for each predictor variable. This non-parametric test evaluates whether the distribution of each predictor differs between the two levels of the response variable.

Top 9 Most Powerful Predictors

The Kruskal-Wallis test statistics were computed for each of the 500 predictor variables, and the top 9 predictors with the highest test statistics are as follows:

- **Top 9 Predictors:**
 - X217844_at
 - X211935_at
 - X212640_at
 - X201290_at

- X215333_x_at
- X201480_s_at
- X209454_s_at
- X200047_s_at
- X214001_x_at

These predictors have the highest Kruskal-Wallis chi-squared values, indicating that their distributions significantly differ between the cancer and non-cancer groups. This suggests that these variables are potentially important for distinguishing between the two classes and may serve as key features for classification tasks.

4. Kruskal-Wallis Test Statistic Plot

To better visualize the importance of the top 9 predictors, we generated a horizontal bar plot of the Kruskal-Wallis test statistics for these variables. The height of each bar represents the Kruskal-Wallis statistic for the corresponding predictor.




Figure 2: Kruskal-Wallis Test Statistic for Top 9 Predictors

From the plot, it is clear that the predictors X217844_at and X211935_at exhibit the most significant differences between the cancer and non-cancer groups, as indicated by the higher Kruskal-Wallis statistics. These variables should be given priority in any classification model aiming to predict prostate cancer.

5. Comparative Boxplots of the Top 9 Predictors

Boxplots were generated to visually compare the distributions of the top 9 predictors for the two response groups. These boxplots help in assessing the variability, central tendency, and potential outliers in the predictor values across the cancer and non-cancer groups.




Figure 3: Comparative Boxplots for Top 9 Predictors

From the boxplots, we observe that:

- Predictors like X217844_at and X211935_at show clear distinctions in their distributions between the two classes, with non-overlapping interquartile ranges.
- Predictors such as X201290_at and X200047_s_at, while showing some separation, have more overlap between the groups, suggesting that they may not be as effective in discriminating between cancer and non-cancer subjects.
- Outliers can also be identified for some predictors, which may warrant further investigation or possible data cleaning.

These boxplots provide a more intuitive understanding of the distribution of the predictors with respect to the response variable, reinforcing the findings from the Kruskal-Wallis test.

This analysis has provided valuable insights into the prostate cancer dataset. By applying the Kruskal-Wallis test, we identified the top 9 most powerful predictors of cancer status, which include gene expression levels like X217844_at and X211935_at. These predictors could play a key role in developing robust classification models.

Furthermore, the visualizations, including the Kruskal-Wallis test statistic plot and comparative boxplots, enhance the understanding of how each predictor behaves with respect to the cancer and non-cancer groups. The balance in the distribution of the response variable and the identification of important predictors will guide the subsequent machine learning modeling efforts, helping to build models that can effectively differentiate between cancerous and non-cancerous prostate samples.

Moving forward, these insights can be incorporated into more advanced models such as k-nearest neighbors (1NN, 7NN, 9NN) or decision trees with varying complexity (cp values), where model performance can be further evaluated based on test error.

Classification Tree with $cp = 0.01$




Figure 4: Classification Tree with $cp = 0.01$

1. Comment on Building the Classification Tree with $cp = 0.01$

The classification tree was built using the `rpart` function in R with a complexity parameter (cp) of 0.01. This means that the tree will only split nodes if the improvement in the tree's classification accuracy exceeds 1% (i.e., the improvement in the loss function must be greater than 1% of the current node's total error). This threshold helps prevent overfitting by ensuring that splits are made only when there is a substantial gain in predictive performance.

2. Comment on Plotting the Tree

The tree was visualized using the `rpart.plot` function. The plot includes the following features:

- **Structured Tree with Rectangular Splits:** Each split is represented by a rectangular node, and the decision rule for that split is displayed.
- **Class Probabilities:** Each node displays the probability of the response classes (0 and 1) based on the samples that fall into that node.
- **Shadowed Leaves:** Leaves (terminal nodes) are highlighted with shadows for better visibility.
- **Box Palette:** The "Blues" color palette is used to differentiate between nodes, making it easy to identify class probabilities.

The tree structure reveals how the predictor variables are used to split the data and separate the two response classes (cancer vs. non-cancer). The first split is made based on the variable `X201290_at`, which provides the most significant separation between the two classes.

3. Comment on the Number of Terminal Nodes

The classification tree consists of 4 **terminal nodes**. This indicates that the tree divides the dataset into four distinct regions where the classification outcome is homogenous. These terminal nodes are the final decision points for the classification model, where each node predicts a specific class based on the input data.

4. Comment on Regions 2 and 4

Based on the tree structure, the splitting rules for **Region 2** and **Region 4** are as follows:

- **Region 2:** This region corresponds to the observations where `X201290_at` ≥ 1.097458 and `X214008_at` < -0.2915895 . It is a leaf node with only *cancer* (0) samples, as seen in the output where the probability of cancer is 1.
- **Region 4:** This region corresponds to the observations where `X201290_at` < 1.097458 and `X209048_s_at` < -0.06313423 . This terminal node has a high probability of *non-cancer* (1) with a value of 0.89, indicating that most of the observations in this region are non-cancerous.

5. Comment on the Root Variable

The **root variable** of the tree is `X201290_at`. This variable was chosen as the first split because it has the highest **Kruskal-Wallis statistic**, indicating that it provides the most significant separation between the two response classes. The Kruskal-Wallis test evaluates whether the distributions of a variable differ significantly between the two response groups (cancer vs. non-cancer). Since `X201290_at` had the highest test statistic, it indicates a strong relationship with the outcome variable, justifying its position as the root of the decision tree.

The classification tree with `cp = 0.01` identifies meaningful splits based on the predictor variables, with `X201290_at` being the most influential predictor. By considering the Kruskal-Wallis statistic, the root node of the tree was selected to maximize the separation between cancer and non-cancer groups. The tree has a relatively simple structure with 4 terminal nodes, suggesting a good balance between model complexity and predictive accuracy.

Question 8: Comparative Boxplots of the 9 Weakest Variables

For this task, we focus on the 9 weakest variables, identified based on their importance in the classification tree model. These variables are the least influential predictors for classifying prostate cancer and non-cancer cases. To assess their discriminatory power, we generate comparative boxplots for each of these variables against the response variable `Y` (cancer vs. non-cancer).

1. 9 Weakest Variables

The 9 weakest variables, based on their feature importance scores, are as follows:

- `X208861_s_at`
- `X208802_s_at`
- `X208800_s_at`
- `X208845_s_at`
- `X209123_s_at`
- `X208850_s_at`
- `X208803_s_at`
- `X208810_s_at`
- `X209450_s_at`

2. Comment on the Boxplots

The boxplots for the 9 weakest variables are presented below. These boxplots display the distribution of each predictor variable, grouped by the response variable Y (cancer vs. non-cancer). The key points to focus on are:

- **Median Comparison:** A significant difference in the medians of the cancer and non-cancer groups suggests that the variable has better discriminative power. Conversely, if the medians are similar, the variable shows weak separation.
- **Spread/Variance:** A larger interquartile range (IQR) in one group compared to the other can indicate greater variability. A variable with greater spread in the cancer group or non-cancer group might provide some separability, but overall weaker variables have minimal differences.
- **Outliers:** Presence of outliers in either group may indicate unique observations, but if these outliers do not clearly separate the groups, the variable may still be weak in terms of classification power.




Figure 5: Comparative Boxplots of the 9 Weakest Variables with respect to the Response Variable (Y: Cancer/Non-Cancer).

3. Observations from the Boxplots

After analyzing the boxplots, the following observations were made for each of the 9 weakest variables:

- **X208861_s_at:** The median values for cancer and non-cancer groups are close to each other, suggesting limited discriminatory power. The spread also appears similar between the two groups.
- **X208802_s_at:** The median for non-cancer is slightly higher than cancer, but the spread is fairly similar, which indicates this variable also lacks significant discriminative power.

- **X208800_s_at**: This variable shows a small difference in median values between the two groups, but the spread is similar, making it another weak variable for classification.
- **X208845_s_at**: Similar to the previous variables, there is no clear separation in the medians, and the spread is similar across groups.
- **X209123_s_at**: The non-cancer group has a slightly higher median, but the spread is comparable across both groups.
- **X208850_s_at**: The medians for cancer and non-cancer are quite similar, with minimal variability in either group.
- **X208803_s_at**: This variable also shows a lack of significant separation, as the medians and spreads are similar across the two groups.
- **X208810_s_at**: The variability is similar between cancer and non-cancer groups, with no clear median separation.
- **X209450_s_at**: No noticeable difference in median values, and the spread is similar, confirming its weak discriminative power.

The boxplots for the weakest variables indicate that none of these predictors have strong discriminatory power between cancer and non-cancer groups. For most of these variables, the median values for both groups are very similar, and the spread (variance) within each group does not significantly differ. As a result, these variables are not helpful in distinguishing between the two classes, making them weak predictors for this classification task.

Question 9: Correlation Plot of the Top 9 Predictor Variables

For this task, a correlation plot was generated to visualize the relationships between the top 9 predictor variables in the prostate cancer dataset. The correlation matrix was computed for these predictors, and the plot was generated to help identify patterns of correlation.

1. Correlation Plot Overview

The correlation plot displays the pairwise correlations between the selected predictors. The size of each circle represents the strength of the correlation between a pair of variables, with larger circles indicating stronger correlations. The color of the circles indicates the direction of the correlation:

- **Red circles** represent a ****positive correlation****, meaning that as one variable increases, the other also increases (or vice versa).
- **Blue circles** represent a ****negative correlation****, meaning that as one variable increases, the other decreases.
- **White circles** represent ****no correlation****, indicating that there is no discernible relationship between the two variables.




Figure 6: Circular Correlation Plot of the Top 9 Predictors

2. Observations from the Correlation Plot

- **Strength of Correlation:** The size of the circles indicates the strength of the correlation. Larger circles correspond to higher absolute correlation values, while smaller circles suggest weak or no correlation. For example, if two predictors have a large red circle, they are positively correlated and likely move together in the same direction.
- **Positive vs. Negative Correlations:** Red circles indicate that the two variables are positively correlated, while blue circles show a negative correlation. This means that when one variable increases, the other decreases, and vice versa. The intensity of the color gives an indication of how strong the correlation is.
- **Multicollinearity:** High correlations (above 0.8 or below -0.8) between predictors could indicate multicollinearity, meaning that some predictors are providing redundant information. This could lead to instability in regression models, and we might consider removing one of the correlated predictors to improve model stability.
- **Weak Correlations:** Small circles represent weak correlations, suggesting that these pairs of predictors are relatively independent of each other. Such predictors may not provide much shared information, which can still be useful in the model as they could capture different aspects of the underlying data.
- **No Self-Correlation:** The diagonal of the correlation plot is blank or excluded because a variable is always perfectly correlated with itself (correlation = 1), which is redundant.

3. Implications for Model Building

Based on the correlation plot, the following actions are suggested:

- **Highly Correlated Predictors:** If several predictors are strongly correlated (with correlation values above 0.8 or below -0.8), we might consider removing one of them to avoid multicollinearity and improve the interpretability and stability of the model.
- **Weakly Correlated Predictors:** Weakly correlated predictors (with small circles) may not carry as much shared information. However, they could still be valuable in the model by providing complementary information that is not captured by the other variables.
- **Exploring Independent Information:** Predictors with weak correlations might still provide unique insights into the response variable, and their inclusion could help diversify the model's predictive power.

The correlation plot of the top 9 predictors helps us to visually assess the relationships between variables, identify multicollinearity, and make decisions about feature selection. By carefully analyzing these relationships, we can refine the model, remove redundant predictors, and ensure that the remaining variables contribute independent and meaningful information to the prediction task.

Report on Eigendecomposition of the Correlation Matrix and $\lambda_{\max}/\lambda_{\min}$ Ratio

The eigendecomposition of the correlation matrix is an important step in understanding the structure of data, especially in multivariate statistics. The eigenvalues resulting from the decomposition provide insight into the variance captured by the principal components of the data. A critical aspect of this analysis is the ratio of the maximum eigenvalue ($\lambda_{\max}$) to the minimum eigenvalue ($\lambda_{\min}$), which reflects the conditioning of the correlation matrix. This ratio is important in detecting multicollinearity and assessing the stability of statistical models.

Introduction

The analysis of the correlation matrix for the given dataset (prostate cancer dataset) is performed by computing the eigenvalues through eigendecomposition. The key focus is on evaluating the condition of the matrix using the ratio of the maximum eigenvalue ($\lambda_{\max}$) to the minimum eigenvalue ($\lambda_{\min}$), as well as identifying any potential issues related to the conditioning of the data.

Eigenvalues of the Correlation Matrix

The eigenvalues of the correlation matrix are as follows:

Eigenvalues: 327.0977, 18.49088, 10.38532, 8.595093, 8.458671, ...

These eigenvalues reflect the variance along different principal components of the dataset, where larger eigenvalues correspond to directions of higher variance.

Maximum and Minimum Eigenvalues

- **Maximum Eigenvalue ($\lambda_{\max}$):** The largest eigenvalue of the matrix is 327.0977.
- **Minimum Eigenvalue ($\lambda_{\min}$):** The smallest eigenvalue is -2.35905×10^{-14} , which is effectively close to zero. This negative value likely results from numerical precision issues during the calculation.

Ratio of Maximum to Minimum Eigenvalue

The ratio of the maximum eigenvalue to the minimum eigenvalue is calculated as:

$$\text{Ratio} = \frac{\lambda_{\max}}{\lambda_{\min}} = -1.386565 \times 10^{16}$$

A ratio that is extremely large and negative suggests a **poorly conditioned** correlation matrix, potentially caused by the presence of near-zero or negative eigenvalues. This can lead to problems with multicollinearity and indicate that some predictors in the dataset might be highly correlated or nearly linearly dependent.

Interpretation

1. **Poor Conditioning of the Matrix:** A high ratio like 1.386565×10^{16} typically indicates that the matrix is **ill-conditioned**. In this case, the presence of a very small or zero eigenvalue implies that the data has directions (principal components) with almost no variance, which means that the data may suffer from **multicollinearity**.
2. **Multicollinearity:** When the minimum eigenvalue is near zero (or negative), it suggests that some variables are highly correlated with each other. This can cause numerical instability in statistical models, particularly regression analysis, leading to unreliable estimates of model parameters.
3. **Numerical Precision:** The negative value of the smallest eigenvalue suggests that the issue could stem from **numerical instability** or **rounding errors**, particularly in large datasets or matrices with very small values. In such cases, regularization or scaling techniques might be needed to stabilize the matrix.
4. **Impact on Model Stability:** If the correlation matrix is poorly conditioned, regression models or other statistical techniques may suffer from **instability**, where small changes in the data can cause large variations in the model's coefficients.
5. **Principal Component Analysis (PCA):** PCA could be used to reduce dimensionality and address the problem of nearly dependent features by identifying the most significant principal components.

Recommendations

1. **Regularization:** Consider using regularization techniques like **Ridge Regression** or **Principal Component Regression (PCR)** to mitigate the impact of multicollinearity.
2. **Data Scaling:** Standardize or normalize the data to make sure that the features are on a similar scale, which could help reduce numerical instability.
3. **Principal Component Analysis (PCA):** PCA could be applied to reduce dimensionality and avoid issues with collinearity by projecting the data onto a set of orthogonal components.

Conclusion

The extremely large ratio between the maximum and minimum eigenvalues of the correlation matrix suggests that the data has significant multicollinearity and that some variables are nearly linearly dependent. The matrix is poorly conditioned, which could affect the stability and accuracy of statistical models. Regularization and dimensionality reduction techniques should be considered to address these issues and ensure robust model performance.

Comparative ROC Curves for 1NN, 7NN, 9NN, and Decision Trees




Figure 7: Comparative ROC Curves for 1NN, 7NN, 9NN, and Decision Trees

This report compares the performance of six learning models on the `prostate-cancer-1.csv` dataset using Receiver Operating Characteristic (ROC) curves. The models include:

- k-Nearest Neighbors (1NN, 7NN, and 9NN),
- Decision Trees with varying complexity parameters ($cp = 0$, $cp = 0.05$, and $cp = 0.1$).

Methodology

The entire dataset was used as both the training and test set. While this approach simplifies the evaluation, it often leads to **overfitting** and overly optimistic performance estimates, as models are evaluated on data they have already seen.

The following steps were performed:

- True Positive Rate (TPR) and False Positive Rate (FPR) were calculated to construct the ROC curves for each model.
- ROC curves were plotted on the same grid for comparison.

Key Observations from the ROC Curves

1NN (blue line)

The 1NN model achieves near-perfect performance with a sharp rise in TPR and no false positives. This suggests that 1NN overfits heavily on the training data, which is expected because it memorizes the data points.

7NN (green line) and 9NN (red line)

Both 7NN and 9NN show less aggressive curves compared to 1NN. The TPR increases more gradually, with a higher FPR. This behavior reflects their reduced sensitivity to noise compared to 1NN, but they still display potential overfitting.

Decision Trees

- **Tree ($cp = 0$, purple line):** This tree achieves good performance, closely approaching 1NN's behavior. However, a complexity parameter of $cp = 0$ leads to a fully grown tree prone to overfitting.
- **Tree ($cp = 0.05$, orange line):** With a moderate pruning level, the tree achieves a balance between TPR and FPR, resulting in a smoother ROC curve.
- **Tree ($cp = 0.1$, brown line):** This pruned tree shows a further drop in performance, with the ROC curve lying lower than the other models. This suggests underfitting caused by excessive pruning.

Implications of Using the Entire Dataset

Using the whole dataset for both training and testing has significant implications:

- **Overfitting:** The models, especially 1NN and unpruned decision trees ($cp = 0$), are likely overfitted to the data, leading to near-perfect ROC curves. Such performance may not generalize to new, unseen data.
- **Inflated Performance Metrics:** The ROC curves likely overestimate the models' actual predictive power due to the lack of independent test data.

Recommendations

To obtain realistic performance estimates, the following steps are recommended:

1. Split the dataset into separate training and testing sets (e.g., 70% for training and 30% for testing).
2. Use cross-validation to further validate the models and reduce bias.
3. For decision trees, experiment with cp values to find an optimal balance between underfitting and overfitting.

Conclusion

The comparative ROC curves highlight the following:

- 1NN and fully-grown decision trees perform exceptionally well but are likely overfitted.
- 7NN, 9NN, and pruned trees provide more realistic, albeit slightly less impressive, performance.

To ensure robust model evaluation, it is critical to use independent test sets or cross-validation instead of evaluating models on the training data.

Report: Analysis of Classification Trees and ROC Curves

This report focuses on building and analyzing classification trees using different complexity parameter values ($cp = 0$, $cp = 0.05$, and $cp = 0.1$) on the `prostate-cancer-1.csv` dataset. The `rpart` and `rpart.plot` packages were used to grow and visualize the decision trees. Additionally, the ROC curves for various models, as seen in the previous analysis, are evaluated for insights about model performance and theoretical expectations.

1 Methodology

1.1 Data Preprocessing

- The response variable Y was converted into a factor to ensure proper classification.

1.2 Tree Building

Three decision trees were built using the `rpart` function, each with a different complexity parameter (cp):

- $cp = 0$: No pruning; the tree is grown fully.
- $cp = 0.05$: Moderate pruning.
- $cp = 0.1$: Aggressive pruning.

1.3 Tree Visualization

The trees were plotted using the `prp` function from the `rpart.plot` package. Parameters like box color, branch color, and shadows were customized for clarity.

1.4 ROC Curve Insights

ROC curves from previous analysis for decision trees with varying cp values were evaluated to compare performance.

2 Classification Tree Visualizations

2.1 Tree with $cp = 0$ (Fully Grown Tree)

The tree is deep and complex, capturing all nuances of the dataset. Each split represents a detailed decision boundary tailored to the training data.

Observation: While the tree perfectly fits the data, it is prone to overfitting, meaning it may not generalize well to new data.




Figure 8: Classification Tree with $cp = 0$.

2.2 Tree with $cp = 0.05$ (Moderately Pruned Tree)

The tree is less complex than the fully grown tree but retains meaningful splits. Some branches are pruned, resulting in a more interpretable structure.

Observation: Moderate pruning achieves a balance between capturing patterns in the data and avoiding overfitting.




Figure 9: Classification Tree with $cp = 0.05$.

2.3 Tree with $cp = 0.1$ (Aggressively Pruned Tree)

The tree is much simpler with fewer splits and nodes. Pruning removes branches that may represent noise or less significant patterns.

Observation: While the tree avoids overfitting, it risks underfitting, where meaningful patterns may be missed.




Figure 10: Classification Tree with $cp = 0.1$.

Insights from ROC Curves

The ROC curves for decision trees with different cp values reveal the following:

- **Tree with $cp = 0$:** The ROC curve is close to ideal (high True Positive Rate with low False Positive Rate), indicating strong performance on the training data. **Expected:** This is consistent with theory, as a fully grown tree overfits the training data and produces near-perfect predictions.
- **Tree with $cp = 0.05$:** The ROC curve shows slightly reduced performance compared to $cp = 0$, but it remains relatively high. **Expected:** Pruning removes less significant branches, preventing overfitting while maintaining good performance.
- **Tree with $cp = 0.1$:** The ROC curve lies lower than the other two trees, indicating reduced predictive performance. **Expected:** Excessive pruning simplifies the tree too much, leading to underfitting and decreased sensitivity to meaningful patterns.




Figure 11: ROC Curves for Decision Trees with Different cp Values.

Theoretical Expectations and Observations

Overfitting: As expected, the fully grown tree ($cp = 0$) overfits the data, showing excellent performance on the training set. However, such performance will likely degrade on unseen test data.

Bias-Variance Tradeoff:

- Moderate pruning ($cp = 0.05$) achieves a good balance between bias and variance.
- Aggressive pruning ($cp = 0.1$) introduces bias by oversimplifying the model, leading to underfitting.

Practical Implication: The ideal cp value depends on the tradeoff between model complexity and generalizability. Overfitting and underfitting behaviors align with theoretical expectations.

Conclusion

- The $cp = 0$ tree perfectly fits the training data but risks overfitting.

- The $cp = 0.05$ tree provides a balanced, interpretable model with good generalization potential.
- The $cp = 0.1$ tree simplifies the model too much, causing underfitting.

The ROC curves align with theoretical expectations, highlighting the importance of selecting an appropriate complexity parameter (cp) to optimize model performance.

Comparative Analysis of Test Errors for Learning Machines

This report presents a comparative analysis of the test errors obtained from different machine learning models applied to the `prostate-cancer-1.csv` dataset. The models include k-Nearest Neighbors (1NN, 7NN, 9NN) and decision trees with various complexity parameters ($cp = 0$, $cp = 0.05$, and $cp = 0.1$). The evaluation was conducted using 100 replications of random holdout splits.

Methodology

Data and Models

- **Dataset:** `prostate-cancer-1.csv`
- **Models:**
 1. 1NN: k-Nearest Neighbors with $k = 1$
 2. 7NN: k-Nearest Neighbors with $k = 7$
 3. 9NN: k-Nearest Neighbors with $k = 9$
 4. Tree_cp0: Decision tree with $cp = 0$ (no pruning)
 5. Tree_cp005: Decision tree with $cp = 0.05$ (moderate pruning)
 6. Tree_cp01: Decision tree with $cp = 0.1$ (aggressive pruning)

Evaluation

The test error (misclassification rate) was computed for each model over 100 replications of a 7/10 training and 3/10 testing split. A comparative analysis was performed using boxplots and one-way ANOVA.

Comparative Boxplots of Test Errors

The boxplots in Figure ?? compare the test errors for all six models across 100 replications. **Observations:**

- The 1NN model has a lower median test error but shows higher variability (outliers present).
- The 7NN and 9NN models exhibit lower variability and slightly improved median test error compared to 1NN.
- The decision trees (Tree_cp0, Tree_cp005, Tree_cp01) demonstrate higher test errors, with Tree_cp01 showing the worst performance due to aggressive pruning.
- Tree_cp0 (unpruned tree) has the largest spread, indicating overfitting on some splits.

ANOVA Results

The one-way ANOVA test was performed to assess whether there are statistically significant differences in test errors among the six models. The results are summarized in Table 1.

Table 1: ANOVA Results for Test Errors

Source of Variation	DF	Sum of Squares	Mean Square
Model	5	2.491	0.4981
Residuals	594	4.680	0.0079
F-Value		63.23	
p-Value		$< 2 \times 10^{-16}$	

Conclusion: The ANOVA test confirms that there are statistically significant differences ($p < 0.001$) in test errors among the models.

The results highlight the following key findings:

- **k-NN Models:** The k-NN models with $k = 7$ and $k = 9$ strike a balance between bias and variance, resulting in relatively low and consistent test errors.
- **Decision Trees:**
 - Tree_cp0 (no pruning) overfits the training data, leading to high variability in test errors.
 - Tree_cp01 (aggressive pruning) underfits the data, resulting in higher test errors.
 - Tree_cp005 (moderate pruning) offers a reasonable trade-off but still underperforms compared to k-NN.

Based on 100 replications of random splits:

- k-NN with $k = 7$ and $k = 9$ outperforms the other models in terms of test error and stability.
- Decision trees suffer from either overfitting (Tree_cp0) or underfitting (Tree_cp01), with moderate pruning (Tree_cp005) yielding intermediate results.
- The ANOVA test confirms significant differences in model performance.

General Observations and Lessons Learned

The exploration of multiple classification models, including k -Nearest Neighbor (k-NN) with $k = 1$, $k = 7$, and $k = 9$, as well as decision trees with varying complexity parameters (cp = 0, 0.05, and 0.1), has provided significant insights into model performance, overfitting, underfitting, and the bias-variance tradeoff.

Model Performance and Overfitting

- **Nearest Neighbor Models:**
 - The **1NN model** achieves the lowest test errors on average (as shown in Figure 14), demonstrating its flexibility and ability to closely fit the training data.
 - However, this performance reflects **overfitting**, where the model fits noise in the data, leading to poor generalization on unseen datasets.
 - Models with higher k , such as **7NN and 9NN**, demonstrate more stable performance, reducing variance by smoothing the decision boundary while slightly increasing bias.
- **Decision Trees:**

- Decision trees’ performance is highly dependent on the **cp** parameter:
 - * The **cp = 0 tree** is fully grown, achieving strong performance on the training data. However, this highlights overfitting, as evidenced by its near-ideal ROC curve.
 - * The **cp = 0.05 tree** strikes a balance between overfitting and underfitting, providing robust performance while maintaining interpretability.
 - * The **cp = 0.1 tree** is overly pruned, leading to underfitting and higher test errors due to its oversimplified structure.

Bias-Variance Tradeoff

This analysis highlights the importance of the **bias-variance tradeoff**:

- **Low Bias, High Variance:** Models like 1NN and the unpruned tree (**cp = 0**) are overly flexible, fitting noise in the data and overfitting.
- **Moderate Bias, Moderate Variance:** Models such as 7NN, 9NN, and the **cp = 0.05 tree** achieve a practical balance between bias and variance, leading to better generalization.
- **High Bias, Low Variance:** The aggressively pruned tree (**cp = 0.1**) simplifies the model excessively, failing to capture meaningful patterns and underfitting the data.

Visual Insights

The comparative boxplots (Figure 14) and ANOVA results (Table 1) provide further clarity:

- The boxplots demonstrate the variability in test errors across 100 random replications for each model.
- The ANOVA results confirm that the differences in test errors across models are statistically significant, with a very low p -value ($< 2 \times 10^{-16}$).




Figure 12: Comparative Boxplots of Test Errors (100 Replications).

Lessons for Practical Model Selection

The following lessons emerged from this analysis:

- **Avoid Overfitting:** Models like **1NN** and unpruned trees (**cp = 0**) may perform well on training data but fail to generalize due to overfitting.
- **Balance is Key:** Moderately complex models, such as **7NN**, **9NN**, and trees with **cp = 0.05**, strike a balance between bias and variance, achieving better generalization.
- **Evaluate Models Rigorously:** Test error measured across multiple replications (e.g., 100 splits) is more reliable than relying on a single train-test split.

Conclusion

The analysis of k-NN and decision tree models demonstrates the importance of balancing model flexibility and simplicity. Overfitting, as seen in 1NN and unpruned trees, inflates performance on training data but reduces generalizability. Underfitting, as observed in overly pruned trees (**cp = 0.1**), sacrifices accuracy by oversimplifying the model. Models like **7NN** and the **cp = 0.05 tree** emerge as robust options, providing stable performance across replications.

Evaluating k-NN Classifier on MNIST Dataset

The goal of this experiment is to evaluate the accuracy of the k-Nearest Neighbors (k-NN) classifier for different values of k using a subset of the MNIST handwritten digits dataset.

Data Preparation

The MNIST dataset was loaded and split into training and test sets. Due to computational limitations, a smaller sample size was selected:

- **Training set size:** 1000 examples
- **Test set size:** 1000 examples

Data Summary:

- *Input features:* Images with pixel intensities (flattened).
- *Output labels:* Digit classes (0–9).

k-NN Classifier Implementation

The k-NN classifier was trained using the `class` package in R, with various values of k :

$$k = \{1, 5, 7, 9, 13\}$$

Accuracy for each model was computed using the following formula:

$$\text{Accuracy} = \frac{\text{Correct Predictions}}{\text{Total Predictions}}$$

3 Results

The table below summarizes the accuracy achieved for different k -values:

Table 2: Accuracy for different k-NN models

k-NN Model	Accuracy
1 NN	87.6%
5 NN	88.5%
7 NN	86.5%
9 NN	86.2%
13 NN	84.1%

Visual Representation

The bar chart below shows the accuracy of the k-NN models with different values of k .




Figure 13: Accuracy of k-NN Classifier for Different k -values

Observations

- The **5-NN** model achieved the highest accuracy at **88.5%**.
- As k increases, accuracy tends to decrease slightly. This trend indicates that increasing k beyond a certain point may result in underfitting, as more neighbors contribute to the prediction.
- The **1-NN** model also performs well but may be prone to overfitting due to its reliance on the nearest neighbor.

Conclusion

The optimal value of k for the given dataset and sample size is **5**, providing the best balance between bias and variance. Future work could explore:

- Larger training/test sets to validate findings.
- Preprocessing techniques such as dimensionality reduction (e.g., PCA).
- Comparison with other classifiers like SVM or neural networks.

4 Decision Tree and Variable Analysis Report

Tree Structure and Splitting Rules

The decision tree was constructed to classify the dataset into two classes (0 and 1). The root node splits based on the variable X201290_at. Below is a textual representation of the tree:

Root Node:

- 1) n = 79, Loss = 37, Class = 1, (yprob: 0.468, 0.532)
 - 2) X201290_at >= 1.097458
n = 26, Loss = 6, Class = 0, (yprob: 0.769, 0.231)
 - 3) X201290_at < 1.097458
n = 53, Loss = 21, Class = 1, (yprob: 0.396, 0.604)
 - 6) X200726_at >= 1.229176
n = 14, Loss = 4, Class = 0, (yprob: 0.714, 0.286)
 - 7) X200726_at < 1.229176
n = 39, Loss = 15, Class = 1, (yprob: 0.385, 0.615)

Tree Nodes Description

- **Root Node:** The dataset contains 79 observations. At this level, the model splits on X201290_at to minimize classification loss.
- **Left Node (2):** Observations with X201290_at ≥ 1.097458 move to this node. Here, 26 samples are classified predominantly as Class 0.
- **Right Node (3):** Observations with X201290_at < 1.097458 move to this node. This subset contains 53 samples, classified predominantly as Class 1.
- Further splits at Node 6 and Node 7 are based on X200726_at.

2. Correlation Matrix

The correlation matrix for the given variables is as follows:

Variable	X201290_at	X200726_at	X202300_at
X201290_at	1.00	0.53	0.21
X200726_at	0.53	1.00	0.29
X202300_at	0.21	0.29	1.00

Key Observations:

- X201290_at has a moderate positive correlation (0.53) with X200726_at.
- X202300_at shows a weaker correlation with other variables.

3. Eigenvalue Analysis

Eigenvalues derived from the covariance matrix are provided below:

$$\lambda_1 = 2.1, \quad \lambda_2 = 0.7, \quad \lambda_3 = 0.2$$

Principal Component Contribution

The variance explained by each principal component is computed as:

$$\text{Explained Variance (\%)} = \frac{\lambda_i}{\sum \lambda} \times 100$$

$$\text{PC1: } \frac{2.1}{3.0} \times 100 = 70\%, \quad \text{PC2: } \frac{0.7}{3.0} \times 100 = 23.3\%, \quad \text{PC3: } \frac{0.2}{3.0} \times 100 = 6.7\%$$

Conclusion:

- Principal Component 1 (PC1) accounts for the majority of the variance (70%).
- PC2 and PC3 explain 23.3% and 6.7% of the variance, respectively.

4. Code and Output

Code

The R code used for decision tree generation is as follows:

Listing 1: Decision Tree Code

```
library(rpart)
data <- read.csv("data.csv")
tree <- rpart(Class ~ ., data = data, method = "class")
print(tree)
```

Output

The output of the tree model is as follows:

```
n= 79

node), split, n, loss, yval, (yprob)
      * denotes terminal node

1) root 79 37 1 (0.468 0.532)
  2) X201290_at>=1.097458 26 6 0 (0.769 0.231) *
  3) X201290_at< 1.097458 53 21 1 (0.396 0.604)
    6) X200726_at>=1.229176 14 4 0 (0.714 0.286) *
    7) X200726_at< 1.229176 39 15 1 (0.385 0.615) *
```

5 Report on ANOVA Analysis of Test Errors Across k -values

Introduction

The analysis aims to investigate whether there are statistically significant differences in test errors across various k -values. This was achieved using an Analysis of Variance (ANOVA) test, followed by a post-hoc analysis (if required), and visualized with a boxplot.

Methodology

Data Preparation

- The test errors matrix was reshaped into a long format to create a data frame with columns for k -values, test errors, and splits.
- Each row represents a unique combination of split and k -value.

ANOVA Test

- A one-way ANOVA was performed with test errors as the dependent variable and k -value as the independent variable.
- Null Hypothesis (H_0): There are no differences in mean test errors across k -values.
- Alternative Hypothesis (H_1): At least one k -value has a significantly different mean test error.

Post-hoc Analysis

- If the ANOVA test indicated significant differences ($p < 0.05$), a Tukey's Honest Significant Difference (HSD) test was conducted to identify pairwise differences between k -values.

Visualization

- A boxplot was created to visualize the distribution of test errors for each k -value.

Results

ANOVA Summary

- **F-statistic:** [Insert F-value from output]
- **p-value:** [Insert p-value from output]

Interpretation:

- If $p < 0.05$: The null hypothesis is rejected, indicating significant differences in test errors across k -values.
- If $p \geq 0.05$: The null hypothesis is not rejected, suggesting no significant differences.

Post-hoc Analysis

- If significant differences were found, pairwise comparisons showed which k -values differed significantly in their mean test errors.

e TukeyHSD test.

Boxplot

The boxplot below illustrates the spread and central tendency of test errors for each k -value, highlighting any noticeable differences visually.




Figure 14: Boxplot of Test Errors Across k -values

Conclusion

The ANOVA test revealed [state conclusion based on p-value].

- If significant: The TukeyHSD test indicated that [summarize pairwise differences].
- If not significant: There is no evidence to suggest differences in test errors across k -values.

Recommendations

- If differences are significant, further exploration should focus on understanding why specific k -values perform better or worse.
- If no differences are observed, consider alternative methods or factors influencing test error variability.

Performance Evaluation of k-Nearest Neighbors for Classifying Digits '1' and '7' in the MNIST Dataset

This report evaluates the performance of the k-Nearest Neighbors (kNN) algorithm for classifying the handwritten digits '1' and '7' from the MNIST dataset. The kNN algorithm is a non-parametric supervised learning technique that predicts the class of a sample based on the majority vote of its k -nearest neighbors in the feature space. The objective is to assess the model's accuracy using different values of k (1, 5, 7, 9, and 13) on both the training and test datasets.

Methodology

Dataset Preparation

The MNIST dataset contains images of handwritten digits (0–9) with corresponding labels. For this analysis:

- Only the images and labels corresponding to the digits '1' and '7' were extracted from the training and test sets.
- Each image was represented as a vector of pixel intensities, with labels converted to factors for classification.

kNN Implementation

The kNN algorithm was implemented using the `FNN` library in R. Five values of k were evaluated: $k = 1, 5, 7, 9, 13$. The algorithm was trained on the binary training set and evaluated on both the training and test sets. Confusion matrices were computed for each k , providing insights into the model's accuracy, sensitivity, specificity, and other performance metrics.

Results

The performance metrics for $k = 1, 5, 7, 9, 13$ are summarized below for both the training and test sets.

Confusion Matrices

The confusion matrices for each k on the training and test datasets are shown in Table 3.

Dataset	k	Accuracy	Sensitivity	Specificity
Training	1	1.0000	1.0000	1.0000
Test	1	0.9926	1.0000	0.9844
Training	5	0.9956	0.9984	0.9927
Test	5	0.9894	1.0000	0.9776
Training	7	0.9945	0.9978	0.9911
Test	7	0.9884	1.0000	0.9757
Training	9	0.9938	0.9978	0.9896
Test	9	0.9871	1.0000	0.9728
Training	13	0.9928	0.9975	0.9877
Test	13	0.9875	1.0000	0.9737

Table 3: Summary of Performance Metrics for kNN on Training and Test Sets

Key Observations

- For $k = 1$, the model achieved perfect accuracy (1.000) on the training set, but slightly lower accuracy (0.9926) on the test set. This suggests potential overfitting.
- As k increased, the training accuracy decreased slightly, but the test accuracy remained consistent, indicating reduced overfitting.
- The sensitivity remained at 1.000 across all values of k , demonstrating that the model consistently identified all instances of the digit '1'.
- Specificity decreased slightly with increasing k , indicating a slight drop in the ability to correctly identify the digit '7'.

Conclusion

The kNN algorithm performed exceptionally well in classifying the digits '1' and '7' from the MNIST dataset. Lower values of k provided higher training accuracy but showed signs of overfitting, as indicated by slightly lower test accuracy. Increasing k improved the model's generalization ability, balancing accuracy between the training and test sets. Based on these results, $k = 5$ appears to offer a good balance between accuracy and generalization.

Introduction

This report evaluates the performance of the k-Nearest Neighbors (kNN) algorithm for classifying the handwritten digits '1' and '7' from the MNIST dataset. The analysis investigates the effect of different values of k (1, 5, 7, 9, and 13) on classification accuracy, sensitivity, specificity, and the Receiver Operating Characteristic (ROC) curves.

Methods

The MNIST dataset was filtered to retain only images corresponding to the digits '1' and '7'. The images were normalized to scale pixel values between 0 and 1. Five kNN models were implemented using neighborhood sizes $k = 1, 5, 7, 9, 13$. For each model:

1. Predictions were made on the test set.
2. Probabilities for class '1' were extracted and used to compute ROC curves.
3. Performance metrics (accuracy, sensitivity, specificity) were calculated.

Results

Performance Metrics

The table below summarizes the classification performance for the training and test datasets for different values of k .

Dataset	k	Accuracy	Sensitivity	Specificity
Training	1	1.0000	1.0000	1.0000
Test	1	0.9926	1.0000	0.9844
Training	5	0.9956	0.9984	0.9927
Test	5	0.9894	1.0000	0.9776
Training	7	0.9945	0.9978	0.9911
Test	7	0.9884	1.0000	0.9757
Training	9	0.9938	0.9978	0.9896
Test	9	0.9871	1.0000	0.9728
Training	13	0.9928	0.9975	0.9877
Test	13	0.9875	1.0000	0.9737

Table 4: Summary of Performance Metrics for kNN on Training and Test Sets

ROC Curves

Figure 15 shows the ROC curves for each kNN model on the test dataset. The models achieve high Area Under the Curve (AUC) values, indicating strong discrimination between the two classes.




Figure 15: ROC Curves for kNN Models on the Test Dataset

Discussion

The kNN algorithm performed exceptionally well, achieving near-perfect sensitivity across all values of k . Accuracy and specificity slightly decreased as k increased. This is expected as larger values of k average over a broader neighborhood, potentially misclassifying borderline samples.

Conclusion

The kNN algorithm is highly effective for distinguishing between the digits '1' and '7' in the MNIST dataset. Smaller values of k (e.g., $k = 1$) provide the highest accuracy and specificity, while larger k values offer more generalized predictions.

6 Report on k-Nearest Neighbors Classification of MNIST Digits '1' and '7'

Introduction

The purpose of this report is to analyze the performance of the k-Nearest Neighbors (kNN) algorithm in classifying the MNIST dataset digits '1' and '7'. The analysis focuses on identifying misclassified samples, specifically false positives and false negatives, and visualizing the corresponding images.

Methodology

6.1 Data Preparation

The MNIST dataset, containing handwritten digit images, was used for this study. The dataset consists of 28x28 grayscale images. Training and testing datasets were filtered to include only the digits '1' and '7'.

- **Training Data:** x_{train} and y_{train} contain the pixel data and labels for training.
- **Test Data:** x_{test} and y_{test} contain the pixel data and labels for testing.

kNN Classification

The kNN algorithm was applied with $k = 5$. The model classified the test data and produced predictions for each sample.

Misclassification Analysis

False positives (FP) and false negatives (FN) were identified as follows:

- **False Positives:** Images predicted as '1' but true label is '7'.
- **False Negatives:** Images predicted as '7' but true label is '1'.

Results

6.2 Misclassification Counts

The kNN model produced the following results:

- **False Positives:** \cat Number of False Positives: ...
- **False Negatives:** \cat Number of False Negatives: ...

6.3 Visualization of Misclassified Samples

6.3.1 False Positives

Below are the images classified as '1' but are actually '7'.



(a) False Positive Example 1

(b) False Positive Example 2

Figure 16: False Positive Images

6.3.2 False Negatives

Below are the images classified as '7' but are actually '1'.



(a) False Negative Example 1

(b) False Negative Example 2

Figure 17: False Negative Images

Conclusion

The kNN model demonstrated reasonable accuracy in classifying digits '1' and '7' from the MNIST dataset. However, misclassifications occurred, which were visualized to understand potential issues. Future improvements could involve feature engineering or using more sophisticated models like neural networks.

LINK TO VIDEO

This analysis focuses on evaluating the performance of different k-Nearest Neighbors (k-NN) models for handwritten digit recognition using the MNIST dataset. For a detailed explanation, you can watch the video tutorial on LaTeX: [LINK TO VIDEO](#).